

Baseline evaluation

All 12 supplied inputs | Recorded live outputs and failures

Outcome

4 automated passes; 8 failures, including 2 execution errors. This is a diagnostic baseline, not a production-quality pass or a human semantic sign-off.

Execution provenance

T1 historical v1; T2-T10 v1 with dynamic Gap Analyzer evidence-ID enum. Failed T2 attempt preserved. T11 and T12 reuse T1.

Capture limitations

T5 JSON escaping was restored from rendered text for a quoted real-time phrase; original rendered text retained. T6 and T8 returned UI errors without structured output. These are failures, not fabricated outputs. No external Langfuse tracing is claimed.

How to inspect

Each T1.md through T12.md contains the exact input, checks and captured JSON output. The raw directory retains machine-readable outputs and the failed T2 attempt. T11 and T12 reuse T1 artifacts. T1 Markdown was deterministically rendered from the captured validated structure.

Interpretation

Strict v2 RelayOps generation is evaluated separately. Its complete catalog and review gate should not be used to claim that the free-text pipeline passed this suite. The source baseline itself permits sparse T11 content; its missing fields must not be copied into the complete production sample.

Reproduction

Import the tested baseline Langflow JSON, connect GROQ_API_KEY and use the supplied official_inputs.txt. Run T1 through T10; evaluate the actual T1 PRD as T11 and its story output as T12. Preserve model, schema/prompt revision, traces and failed attempts. Repeated runs can vary.

Cases T1 and T2

Exact inputs, check outcomes and observed response summaries

T1 - PASS

The user should be able to filter reports by date range, category, and status. Results must load in under 2 seconds. PM: Sarah. Deadline: Q3.

PASS: Pipeline did not error; PASS: Source quotes, typed output, references and coverage validate; PASS: All three filters; PASS: Exact 2 second target; PASS: PM Sarah captured as stakeholder; PASS: Q3 deadline

Status draft_ready. Extracted 2 requirements; extraction label sufficient. Reuses the previously captured live T1 run; Markdown was deterministically reconstructed from validated structured output.

Full recorded output: T1.md and raw/T1.output.json. Human review: pending.

T2 - FAIL

We need better reporting. Something like what Competitor X has.

PASS: Pipeline did not error; PASS: Source quotes, typed output, references and coverage validate; PASS: Ambiguous intent; PASS: No invented actionable requirements; FAIL: Metrics, format, users clarified

Status clarification_required. Extracted 0 requirements; extraction label ambiguous. Captured from the live Langflow Playground.

Full recorded output: T2.md and raw/T2.output.json. Human review: pending.

Cases T3 and T4

Exact inputs, check outcomes and observed response summaries

T3 - FAIL

The dashboard should auto-refresh every 5 seconds. [Later in same meeting] Performance is critical, minimize API calls.

PASS: Pipeline did not error; PASS: Source quotes, typed output, references and coverage validate; FAIL: Refresh/call tension flagged; PASS: Stakeholder clarification

Status draft_ready. Extracted 2 requirements; extraction label sufficient. Captured from the live Langflow Playground.

Full recorded output: T3.md and raw/T3.output.json. Human review: pending.

T4 - FAIL

Users need to export reports as PDF and CSV. PDF must include company logo. CSV must preserve formulas.

PASS: Pipeline did not error; PASS: Source quotes, typed output, references and coverage validate; FAIL: Exact required criteria and no extra criteria; FAIL: Both criteria reach stories

Status draft_ready. Extracted 3 requirements; extraction label sufficient. Captured from the live Langflow Playground.

Full recorded output: T4.md and raw/T4.output.json. Human review: pending.

Cases T5 and T6

Exact inputs, check outcomes and observed response summaries

T5 - FAIL

Discussed dashboard... John mentioned something about real-time... budget TBD...

PASS: Pipeline did not error; PASS: Source quotes, typed output, references and coverage validate; FAIL: Insufficient;
PASS: No invented requirements; PASS: Dashboard, real-time, budget gaps

Status clarification_required. Extracted 0 requirements; extraction label ambiguous. Captured from the live Langflow Playground.

Full recorded output: T5.md and raw/T5.output.json. Human review: pending.

T6 - FAIL

Engineering wants microservices. Design wants single-page app. PM wants it shipped by March.

FAIL: Pipeline did not error

Status error. No structured extraction returned. Captured from the live Langflow Playground.

Full recorded output: T6.md and raw/T6.output.json. Human review: pending.

Cases T7 and T8

Exact inputs, check outcomes and observed response summaries

T7 - FAIL

API must support 10,000 concurrent users. Response time < 200ms at p95. Must integrate with Salesforce REST API v52.

PASS: Pipeline did not error; PASS: Source quotes, typed output, references and coverage validate; FAIL: Exact capacity, latency, percentile, API version as NFRs

Status draft_ready. Extracted 3 requirements; extraction label sufficient. Captured from the live Langflow Playground.

Full recorded output: T7.md and raw/T7.output.json. Human review: pending.

T8 - FAIL

Admins need bulk user management. End users need a simplified view. Auditors need read-only access with full history.

FAIL: Pipeline did not error

Status error. No structured extraction returned. Captured from the live Langflow Playground.

Full recorded output: T8.md and raw/T8.output.json. Human review: pending.

Cases T9 and T10

Exact inputs, check outcomes and observed response summaries

T9 - PASS

Meeting happened. Notes: none.

PASS: Pipeline did not error; PASS: Source quotes, typed output, references and coverage validate; PASS: No requirements extractable; PASS: No PRD or stories

Status clarification_required. Extracted 0 requirements; extraction label empty. Captured from the live Langflow Playground.

Full recorded output: T9.md and raw/T9.output.json. Human review: pending.

T10 - FAIL

SSO login requires the new auth service which is being built by Team Alpha. ETA unknown.

PASS: Pipeline did not error; PASS: Source quotes, typed output, references and coverage validate; PASS: SSO extracted; FAIL: Dependency and owner; PASS: Unknown ETA risk

Status draft_ready. Extracted 1 requirements; extraction label sufficient. Captured from the live Langflow Playground.

Full recorded output: T10.md and raw/T10.output.json. Human review: pending.

Cases T11 and T12

Exact inputs, check outcomes and observed response summaries

T11 - PASS

[Use the extraction output from T1 to generate a full PRD]

PASS: Pipeline did not error; PASS: Source quotes, typed output, references and coverage validate; PASS: All ten template headings; PASS: PRD uses only and all T1 evidence

Status draft_ready. Extracted 2 requirements; extraction label sufficient. Evaluates the actual T1 pipeline artifact, deterministically rendered from captured structured output; not an independent model call.

Full recorded output: T11.md and raw/T11.output.json. Human review: pending.

T12 - PASS

[Use the PRD from T11 to generate user stories]

PASS: Pipeline did not error; PASS: Source quotes, typed output, references and coverage validate; PASS: Epics and features; PASS: Standard stories and priorities

Status draft_ready. Extracted 2 requirements; extraction label sufficient. Evaluates the actual T1 pipeline artifact, deterministically rendered from captured structured output; not an independent model call.

Full recorded output: T12.md and raw/T12.output.json. Human review: pending.